

Transaction Management- Intro

Kalyani Selvarajah
School of Computer Science
University of Windsor



Advanced Database Topics
COMP 8157 01
Fall 2023

TODAY'S AGENDA

Transaction

Consistency of Database

Concurrency control

Demo: SQL server



<https://domains.upperlink.ng/elementor-947/>

PRE-ASSESSMENT


1. What do you mean by **Transaction**?
2. What are the potential problems of **Concurrency control**?
3. How to run a transaction in SQL server?



Introduction: Transaction

A [transaction](#) is an action, or series of actions, carried out by user or application, which reads or updates contents of database.

It is the basic unit of change in a DBMS.



Source: <https://www.pinterest.com/pin/775533998318621246/>

TRANSACTION EXAMPLE (1)

Transfer \$200 from Tom's account to Ann's account.

Transaction:

→ Deduct \$200 from Tom's account.

→ Add \$200 to Ann's account.



TRANSACTION EXAMPLE (2)

Staff (staffNo, fName, lName, position, sex, DOB, salary)

PropertyForRent (propertyNo, street, city, postcode, type, rooms, rent, staffNo)



Database operations

Non-database operations

TRANSACTION EXAMPLE (2)

Staff (staffNo, fName, lName, position, sex, DOB, salary)

PropertyForRent (propertyNo, street, city, postcode, type, rooms, rent, staffNo)



I want to delete a staff ID x because he resigned.

DELETE(staffNo = x)

Staff

StaffNo	fName	lName	Position	Sex	DOB	Salary
X	John	White	Manager	M	1-Oct-45	30000
Y	Ann	Beech	Assistant	F	10-Nov-60	12000
Z	David	Ford	Supervisor	M	24-Mar-58	18000

PropertyForRent

PropertyNo	Street	City	Postcode	Type	Rooms	Rent	StaffNo
PA14	16 Holhead	Aberdeen	AB7 5SU	House	6	SA9	Z
PL94	6 Argyll St	London	NW2	Flat	4	SL41	X

TRANSACTION EXAMPLE (2)

Staff (staffNo, fName, lName, position, sex, DOB, salary)

PropertyForRent (propertyNo, street, city, postcode, type, rooms, rent, staffNo)

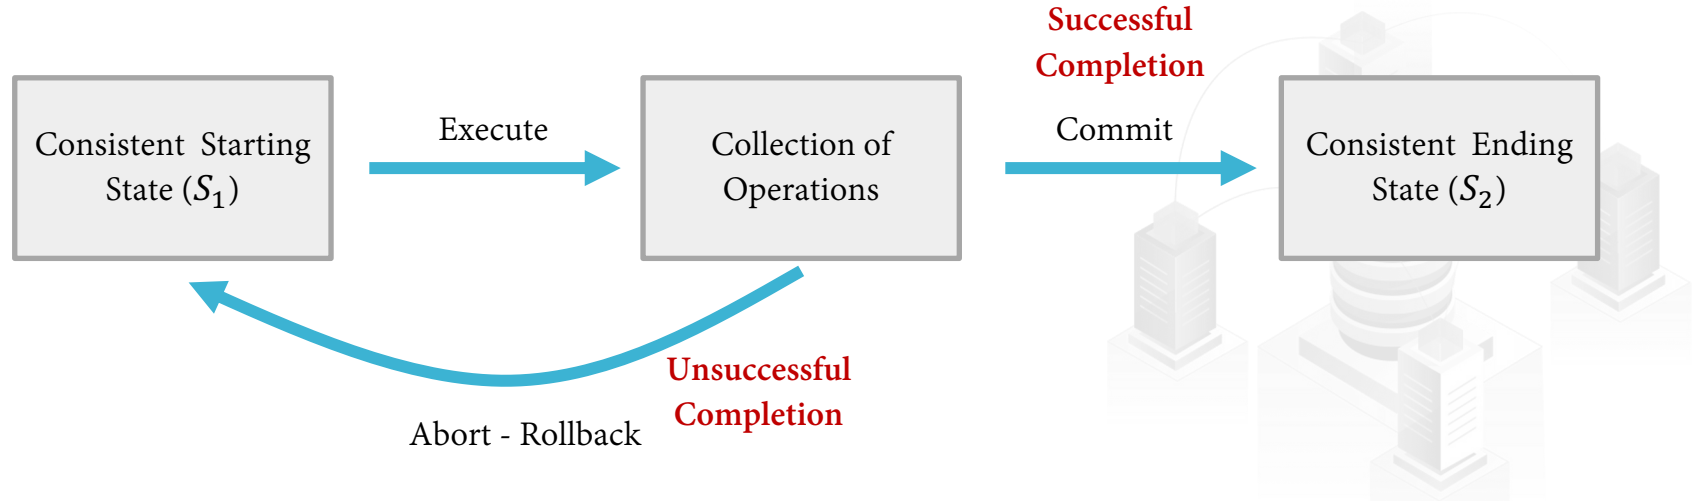


I want to delete a staff ID
x because he resigned.

DELETE(staffNo = x)
FOR ALL PropertyForRent records, pno
BEGIN
READ(propertyNo = pno, staffNo)
IF (staffNo = x) THEN
BEGIN
staffNo = newStaffNo
WRITE(property No = pno, staffNo)
END
END

TRANSACTION LIFE CYCLE

By Jim Gray

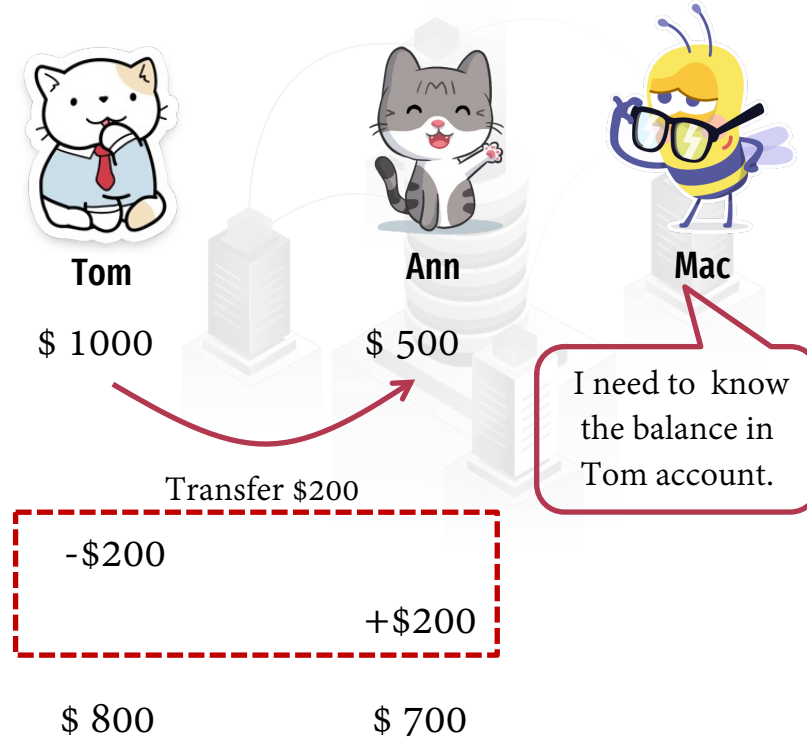
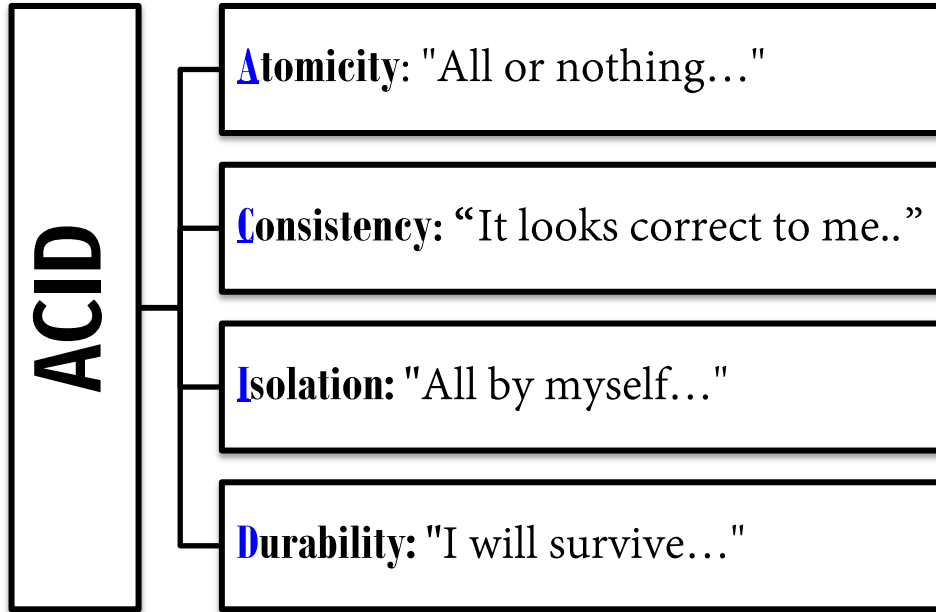


Success - transaction commits and database reaches a new consistent state.

Failure - transaction aborts, and database must be restored to consistent state before it started. Such a transaction is rolled back or undone.


NOTE: If we committed a transaction, it cannot be aborted. However, an aborted transaction that is rolled back can be restarted later.

PROPERTIES OF TRANSACTIONS (ACID)





A




B



CONCURRENCY CONTROL

Concurrent access of data X:



Time	T_1	T_2
t_1		BEGIN
t_2	BEGIN	READ(X)
t_3	READ(X)	$X = X + 100$
t_4	$X = X - 10$	WRITE(X)
t_5	WRITE(X)	COMMIT
t_6	COMMIT	

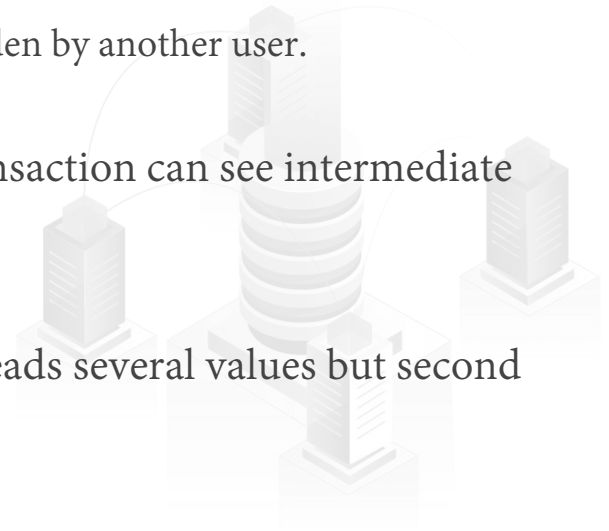


A major objective in developing a database is to **enable many users to access shared data** concurrently.

Concurrency Control is the protocol to allow transactions to access a database in a multi-programmed fashion while preserving the illusion that each of them is executing alone on a dedicated system.

POTENTIAL PROBLEMS OF CONCURRENCY

1. **Lost update problem:** Successfully completed update is overridden by another user.
2. **Uncommitted dependency problem:** Occurs when one transaction can see intermediate results of another transaction before it has committed.
3. **Inconsistent analysis problem:** Occurs when transaction reads several values but second transaction updates some of them during execution of first.
4. **Nonrepeatable read:** Occur when a transaction rereads a data item it has previously read but, in between, another transaction has modified it.




1. LOST UPDATE PROBLEM

Assume initially account X has \$100.

T_1 withdraws \$10 from account X .

T_2 deposits \$100 into same account X .



Time	T_1	T_2	X
t_1		BEGIN	100
t_2	BEGIN	READ(X)	100
t_3	READ(X)	$X = X + 100$	100
t_4	$X = X - 10$	WRITE(X)	200
t_5	WRITE(X)	COMMIT	90
t_6	COMMIT		90

The lost update problem.




2. UNCOMMITTED DEPENDENCY PROBLEM (DIRTY READ)

Assume initially account X has \$100.

T_1 withdraws \$10 from account X .

T_2 deposits \$100 into same account X .



Time	T_1	T_2	X
t_1		BEGIN	100
t_2		READ(X)	100
t_3		$X = X + 100$	100
t_4	BEGIN	WRITE(X)	200
t_5	READ(X)	...	200
t_6	$X = X - 10$	ROLLBACK	100
t_7	WRITE(X)		190
t_8	COMMIT		190

The uncommitted dependency problem.



3. INCONSISTENT ANALYSIS PROBLEM

Assume initially


→ account X has \$100.

→ account Y has \$50.

→ account Z has \$25.

T_1 transfers \$10 from account X to Z .

T_2 analyzes the total of accounts X , Y & Z .



Time	T_1	T_2	X	Y	Z	SUM
t_1		BEGIN	100	50	25	
t_2	BEGIN	SUM=0	100	50	25	0
t_3	READ(X)	READ(X)	100	50	25	0
t_4	X = X-10	SUM = SUM+X	100	50	25	100
t_5	WRITE(X)	READ(Y)	90	50	25	100
t_6	READ(Z)	SUM = SUM+Y	90	50	25	150
t_7	Z = Z+10		90	50	25	150
t_8	WRITE(Z)		90	50	35	150
t_9	COMMIT	READ(Z)	90	50	35	150
t_{10}		SUM = SUM+Z	90	50	35	185
t_{11}		COMMIT	90	50	35	185


The inconsistent analysis problem.

4. NONREPEATABLE READ (PHANTOM READ)

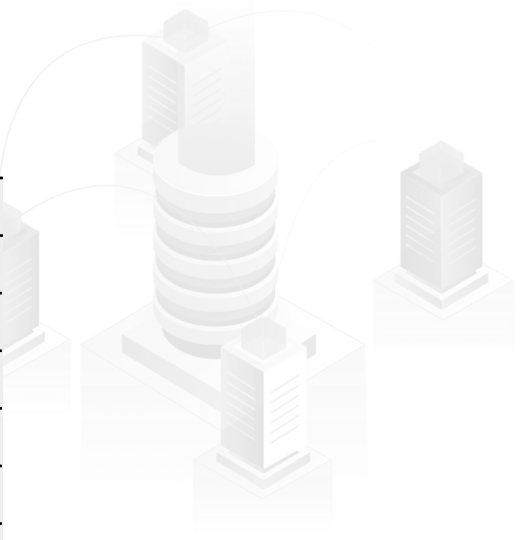
Assume initially account X has \$100.

T_1 withdraws \$10 from account X .

T_2 deposits \$100 into same account X .



Time	T_1	T_2	X
t_1		BEGIN	100
t_2	BEGIN	READ(X)	100
t_3	READ(X)	100
t_4	$X = X - 10$		90
t_5	WRITE(X)		90
t_6	COMMIT	90
		READ (X)	90
		COMMIT	



The unrepeatable read problem .

Serializability and Recoverability

Objective of a concurrency control protocol is to **schedule transactions** in such a way as to avoid any interference.

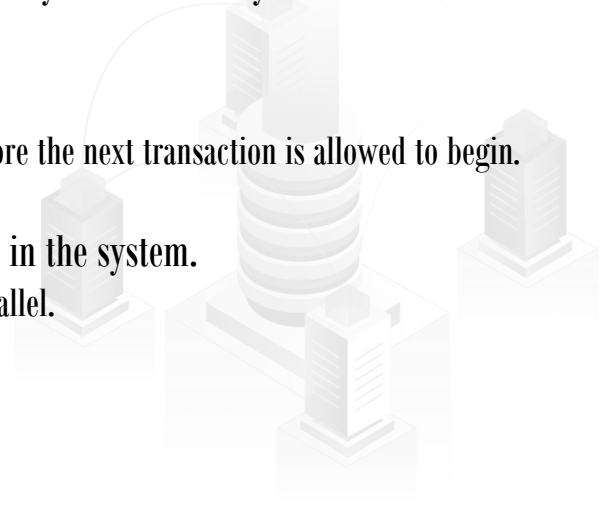
Obvious solution?

allow only one transaction to execute at a time: one transaction is committed before the next transaction is allowed to begin.

Aim of a multi-user DBMS: to maximize the degree of **concurrency or parallelism** in the system.

transactions that can execute without interfering with one another can run in parallel.

How?





Schedule

Schedule: A sequence of the operations by a set of concurrent transactions that preserves the order of the operations in each of the individual transactions.

Serial Schedule: A schedule where the operations of each transaction are executed consecutively without any interleaved operations from other transactions.

Non serial schedule: A schedule where the operations from a set of concurrent transactions are interleaved.



Conflicts

A conflict occurs when two running transactions perform **noncompatible operations** on the same data item of the database.

A conflict occurs when one transaction writes an item that another transaction is reading or writing.

	READ T_2	WRITE T_2
READ T_1	No Conflict	Conflict
WRITE T_1	Conflict	Conflict

CONFLICT MATRIX.

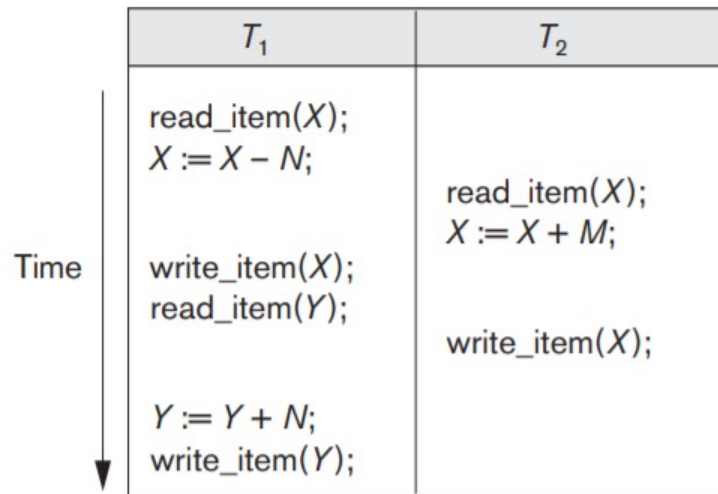
Conflicts: Example

Conflict operations:

- ✓ the operations $r_1(X)$ and $w_2(X)$,
- ✓ the operations $r_2(X)$ and $w_1(X)$,
- ✓ and the operations $w_1(X)$ and $w_2(X)$

Do not Conflict:

- ✓ the operations $r_1(X)$ and $r_2(X)$ -- since they are both read operations.
- ✓ the operations $w_2(X)$ and $w_1(Y)$ -- they operate on distinct data items X and Y .
- ✓ operations $r_1(X)$ and $w_1(X)$ -- they belong to the same transaction.

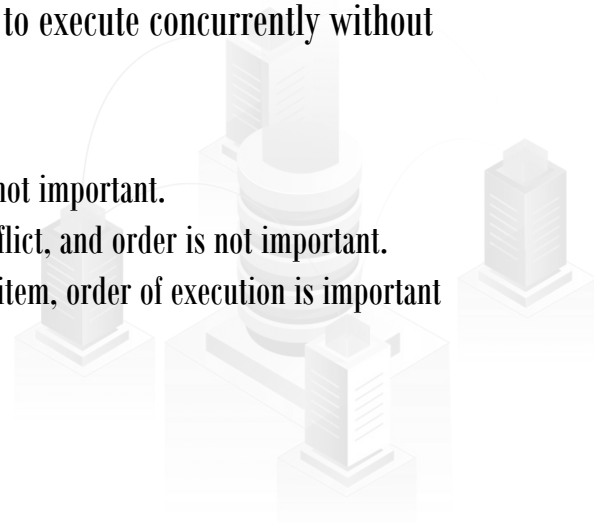


Serializability

Objective of serializability is to find non-serial schedules that allow transactions to execute concurrently without interfering with one another.

In serializability, **ordering of read/writes** is important:

- ✓ If two transactions only read a data item, they do not conflict, and order is not important.
- ✓ If two transactions either read or write separate data items, they do not conflict, and order is not important.
- ✓ If one transaction writes a data item and another reads or writes same data item, order of execution is important



Serializability

TIME

Time	T_1	T_2
t_1	BEGIN	
t_2	READ(X)	
t_3	WRITE(X)	
t_4		BEGIN
t_5		READ(X)
t_6		WRITE(X)
t_7	READ(Y)	
t_8	WRITE(Y)	
t_9	COMMIT	
t_{10}		READ(Y)
t_{11}		WRITE(Y)
t_{12}		COMMIT

(a) Schedule S_1

T_1	T_2
BEGIN	
READ(X)	
WRITE(X)	
	BEGIN
	READ(X)
	WRITE(X)
READ(Y)	
WRITE(Y)	
COMMIT	
	READ(Y)
	WRITE(Y)
	COMMIT

(b) Schedule S_2

T_1	T_2
BEGIN	
READ(X)	
WRITE(X)	
READ(Y)	
WRITE(Y)	
COMMIT	
	BEGIN
	READ(X)
	WRITE(X)
	READ(Y)
	WRITE(Y)
	COMMIT

(c) Schedule S_3

Equivalent schedules: (a) nonserial schedule S_1 ; (b) nonserial schedule S_2 equivalent to S_1 ; (c) serial schedule S_3 , equivalent to S_1 and S_2 .

Conflict serializability

Conflict serializable schedule orders any **conflicting operations** in same way as some serial execution.

Testing for conflict serializability:

Under constrained write rule (transaction updates data item based on its old value, which is first read), use **precedence graph** to test for serializability.

Precedence Graph:

Create:

node for each transaction;


- ✓ a directed edge $T_i \rightarrow T_j$, if T_j reads the value of an item written by T_i ;
- ✓ a directed edge $T_i \rightarrow T_j$, if T_j writes a value into an item after it has been read by T_i .
- ✓ a directed edge $T_i \rightarrow T_j$, if T_j writes a value into an item after it has been written by T_i .

$R_i(A) \rightarrow W_j(A)$

$W_i(A) \rightarrow R_j(A)$

$W_i(A) \rightarrow W_j(A)$


If there is a cycle in the precedence graph, schedule S is not (conflict) serializable; if there is no cycle, S is serializable.



Example - Non-conflict serializable schedule

T_1 is transferring \$100 from one account with balance X to another account with balance Y .

T_2 is increasing balance of these two accounts by 10%.



Precedence graph has a cycle and so is not serializable.




Time	T_1	T_2
t_1	BEGIN	
t_2	READ(X)	
t_3	$X = X + 100$	
t_4	WRITE(X)	BEGIN
t_5		READ(X)
t_6		$X = X * 1.1$
t_7		WRITE(X)
t_8		READ(Y)
t_9		$Y = Y * 1.1$
t_{10}		WRITE(Y)
t_{11}	READ(Y)	COMMIT
t_{12}	$Y = Y - 100$	
t_{13}	WRITE(Y)	
t_{14}	COMMIT	

Two concurrent update transactions that are not conflict serializable.

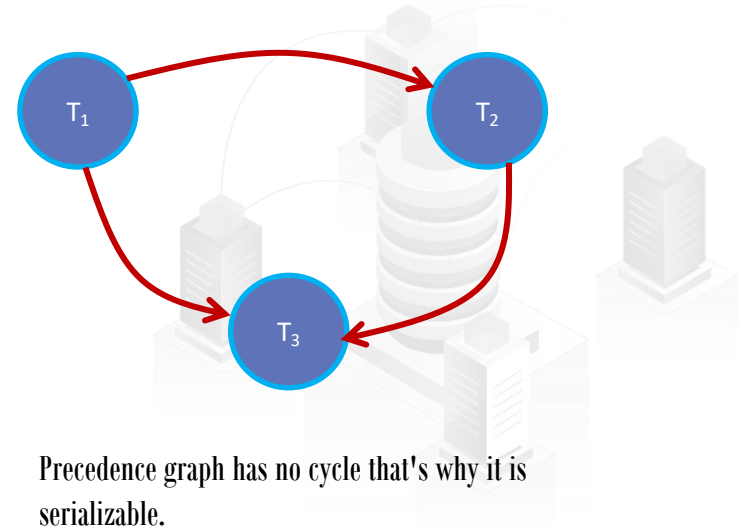
Example

Time	T ₁	T ₂	T ₃
t ₁	READ(A)		
t ₂	WRITE(A)		
t ₃			READ(A)
t ₄			WRITE(A)
t ₅			COMMIT
t ₆		READ(B)	
t ₇		WRITE(B)	
t ₈		COMMIT	
t ₉	READ(B)		
t ₁₀	WRITE(B)		
t ₁₀	COMMIT		



Example

Time	T ₁	T ₂	T ₃
T ₁	READ(A)		
T ₂	READ(C)		
T ₃	WRITE(A)		
T ₄		READ(B)	
T ₅	WRITE (C)		
T ₆		READ(A)	
T ₇			READ (C)
T ₈		WRITE(B)	
T ₉			READ (B)
T ₁₀			WRITE(C)
T ₁₁		WRITE(A)	
T ₁₂			WRITE(B)



View Serializability

There are several other types of serializability that offer less stringent definitions of schedule equivalence than that offered by conflict serializability.

One less restrictive definition is called **view serializability**.


Two schedules S_1 and S_2 consisting of the same operations from n transactions T_1, T_2, \dots, T_n are view equivalent if:

- ✓ For each data item x , if T_i reads initial value of x in S_1 , T_i must also read initial value of x in S_2 .
- ✓ For each read on x by T_i in S_1 , if value read by T_i is written by T_j , T_i must also read value of x produced by T_j in S_2 .
- ✓ For each data item x , if last write on x performed by T_i in S_1 , same transaction must perform final write on x in S_2 .

View Serializability

Time	T ₁	T ₂	T ₃
t ₁	BEGIN		
t ₂	READ(X)		
t ₃			
t ₄		BEGIN	
t ₅		WRITE(X)	
t ₆	WRITE (X)	COMMIT	
t ₇	COMMIT		
t ₈			BEGIN
t ₉			WRITE(X)
t ₁₀			COMMIT

It is view serializable, as it is equivalent to the serial schedule T₁ followed by T₂ followed by T₃



we can see that the edge T₂ → T₁ should not have been inserted into the graph, as the values of x written by T₁ and T₂ were never used by any other transaction because of the blind writes.

S1 = R1(X), W2(X), W1(X), W3(X)
S2 = R1(X), W1(X), W2(X), W3(X)

Schedules S1 and S2 are view equivalent